



THE UNIVERSITY OF TEXAS AT AUSTIN
McCOMBS SCHOOL OF BUSINESS

Model building: dummy variables

Lecture 10

STA 371G

Let's predict fuel economy (miles per gallon) for different car models of the 70s.



Let's predict fuel economy (miles per gallon) for different car models of the 70s.



- Cylinders
- Displacement
- Horsepower
- Weight
- Acceleration
- Year (After 1975 or not)

Exploring the data

```
head(cars)
```

	MPG	Cylinders	Displacement	HP	Weight	Acceleration	After1975	Origin
1	18	8	307	130	3504	12.0	No	US
2	15	8	350	165	3693	11.5	No	US
3	18	8	318	150	3436	11.0	No	US
4	16	8	304	150	3433	12.0	No	US
5	17	8	302	140	3449	10.5	No	US
6	15	8	429	198	4341	10.0	No	US

Exploring the data

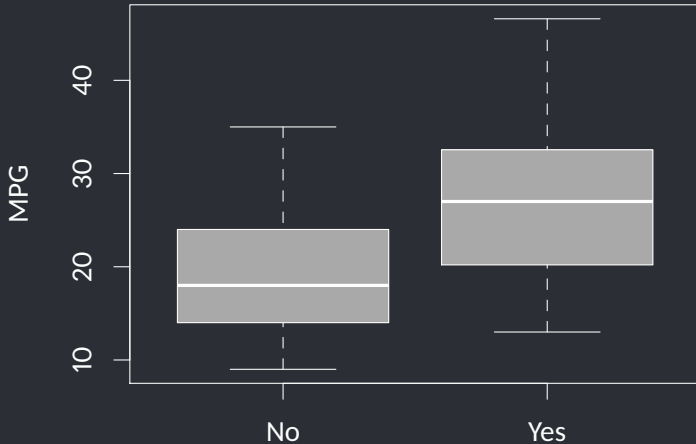
```
head(cars)
```

	MPG	Cylinders	Displacement	HP	Weight	Acceleration	After1975	Origin
1	18	8	307	130	3504	12.0	No	US
2	15	8	350	165	3693	11.5	No	US
3	18	8	318	150	3436	11.0	No	US
4	16	8	304	150	3433	12.0	No	US
5	17	8	302	140	3449	10.5	No	US
6	15	8	429	198	4341	10.0	No	US

How do we handle the Yes/No data in the After1975 column?

Are late-model cars different?

```
boxplot(MPG ~ After1975, data=cars, ylab="MPG",  
        xlab="After 1975", col='darkgray')
```



Exploring the data

To incorporate the After1975 variable into a regression model, we create a **dummy variable** called LateModel that maps a “Yes” to 1, and “No” to 0.

Exploring the data

To incorporate the After1975 variable into a regression model, we create a **dummy variable** called LateModel that maps a “Yes” to 1, and “No” to 0.

```
cars$LateModel <-  
  ifelse(cars$After1975 == "Yes", 1, 0)
```


Exploring the data

To incorporate the After1975 variable into a regression model, we create a **dummy variable** called LateModel that maps a “Yes” to 1, and “No” to 0.

```
cars$LateModel <-  
  ifelse(cars$After1975 == "Yes", 1, 0)
```

Now let's a regression model using the predictors Cylinders, Displacement, HP, Weight, Acceleration, and LateModel.

R will actually create this “dummy” (0/1) variable for us automatically, when you put a categorical variable (what R calls a “factor” into a model).

```
summary(lm(MPG ~ Cylinders + Displacement + HP + Weight +  
Acceleration + After1975, data=cars))
```

Call:

```
lm(formula = MPG ~ Cylinders + Displacement + HP + Weight + Acceleration  
After1975, data = cars)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-8.9302	-2.5727	-0.2574	2.0630	15.0381

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	42.1939988	2.3687735	17.813	< 2e-16	***
Cylinders	-0.5840362	0.3601220	-1.622	0.106	
Displacement	0.0074811	0.0079791	0.938	0.349	
HP	-0.0198909	0.0147848	-1.345	0.179	
Weight	-0.0059904	0.0007194	-8.327	1.46e-15	***
Acceleration	0.0354327	0.1103506	0.321	0.748	
After1975Yes	4.3590043	0.4016220	10.853	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.721 on 385 degrees of freedom

Dummy variables

After1975Yes is 1 whenever After1975 is “Yes,” and 0 otherwise:

MPG	...	Acceleration	After1975	After1975Yes
...
25	...	13.5	No	0
33	...	17.5	No	0
28	...	15.5	Yes	1
25	...	16.9	Yes	1
...

Dummy variables

After1975Yes is 1 whenever After1975 is “Yes,” and 0 otherwise:

MPG	...	Acceleration	After1975	After1975Yes
...
25	...	13.5	No	0
33	...	17.5	No	0
28	...	15.5	Yes	1
25	...	16.9	Yes	1
...

Notice that we do not have a After1975No variable; it would cause problems because it would be perfectly correlated with After1975Yes.

Interpretation of the $\hat{\beta}$ of the dummy variable

Our regression equation is:

$$\widehat{\text{MPG}} = 41.71 - 0.02 \cdot \text{HP} - 0.01 \cdot \text{Weight} + 4.33 \cdot \text{After1975Yes}.$$



Interpretation of the $\hat{\beta}$ of the dummy variable

Our regression equation is:

$$\widehat{\text{MPG}} = 41.71 - 0.02 \cdot \text{HP} - 0.01 \cdot \text{Weight} + 4.33 \cdot \text{After1975Yes}.$$

Let's interpret the coefficient 4.33. Consider this:

- Model A and B have the same HP and Weight.



Interpretation of the $\hat{\beta}$ of the dummy variable

Our regression equation is:

$$\widehat{\text{MPG}} = 41.71 - 0.02 \cdot \text{HP} - 0.01 \cdot \text{Weight} + 4.33 \cdot \text{After1975Yes}.$$

Let's interpret the coefficient 4.33. Consider this:

- Model A and B have the same HP and Weight.
- Model A was manufactured before 1975, whereas B was manufactured after 1975.



Interpretation of the $\hat{\beta}$ of the dummy variable

Our regression equation is:

$$\widehat{\text{MPG}} = 41.71 - 0.02 \cdot \text{HP} - 0.01 \cdot \text{Weight} + 4.33 \cdot \text{After1975Yes}.$$

Let's interpret the coefficient 4.33. Consider this:

- Model A and B have the same HP and Weight.
- Model A was manufactured before 1975, whereas B was manufactured after 1975.
- We predict Model B will have a MPG that is 4.33 higher than Model A.



Interpretation of the $\hat{\beta}$ of the dummy variable

R has assigned “Yes” to 1 and “No” to 0 in our dummy variable, so the “reference level” is cars manufactured before 1975.

Interpretation of the $\hat{\beta}$ of the dummy variable

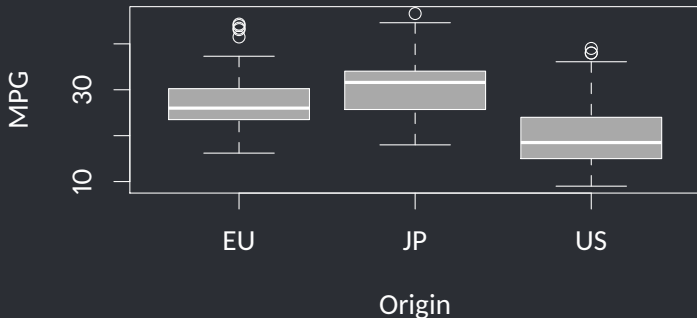
R has assigned “Yes” to 1 and “No” to 0 in our dummy variable, so the “reference level” is cars manufactured before 1975.

If we created a dummy variable `After1975No` that is 1 for cars manufactured *before* 1975, what would the regression look like?

What if there are more than two categories?

The `Origin` variable represents the country of manufacture:

```
boxplot(MPG ~ Origin, data=cars, ylab="MPG",  
        xlab="Origin", col='darkgray')
```



What if there are more than two categories?

- The `Origin` variable has 3 “levels”—US, EU, and JP—so we can’t easily convert this into a 0/1 dummy variable.

What if there are more than two categories?

- The `Origin` variable has 3 “levels”—US, EU, and JP—so we can’t easily convert this into a 0/1 dummy variable.
- The solution is to create a dummy variable for *each* level (category), and include **all but one** of them as predictors in the model.

What if there are more than two categories?

- The `Origin` variable has 3 “levels”—US, EU, and JP—so we can’t easily convert this into a 0/1 dummy variable.
- The solution is to create a dummy variable for *each* level (category), and include **all but one** of them as predictors in the model.
- The category left out is the **reference level** and all slope coefficients for dummy variables are interpreted as the difference between that category and the reference level.

What if there are more than two categories?

- The `Origin` variable has 3 “levels”—US, EU, and JP—so we can’t easily convert this into a 0/1 dummy variable.
- The solution is to create a dummy variable for *each* level (category), and include **all but one** of them as predictors in the model.
- The category left out is the **reference level** and all slope coefficients for dummy variables are interpreted as the difference between that category and the reference level.
- If *any* of the dummy variables are significant for a particular categorical variable, we consider the entire categorical variable to be significant!


```
origin.model <- lm(MPG ~ HP + Weight + After1975 + Origin, data=cars)
summary(origin.model)
```

Call:

```
lm(formula = MPG ~ HP + Weight + After1975 + Origin, data = cars)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-9.705	-2.243	-0.199	1.816	13.789

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	40.181974	0.874271	45.96	<2e-16	***
HP	-0.027984	0.009865	-2.84	0.0048	**
Weight	-0.005160	0.000477	-10.81	<2e-16	***
After1975Yes	4.334280	0.392849	11.03	<2e-16	***
OriginJP	1.000586	0.611954	1.64	0.1029	
OriginUS	-1.592573	0.561900	-2.83	0.0048	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.6 on 386 degrees of freedom

Multiple R-squared: 0.787, Adjusted R-squared: 0.784

F-statistic: 285 on 5 and 386 DF, p-value: <2e-16

A warning about categorical variables with numeric representations

- In the original dataset, the origin was represented as 1 for US, 2 for EU and 3 for JP.

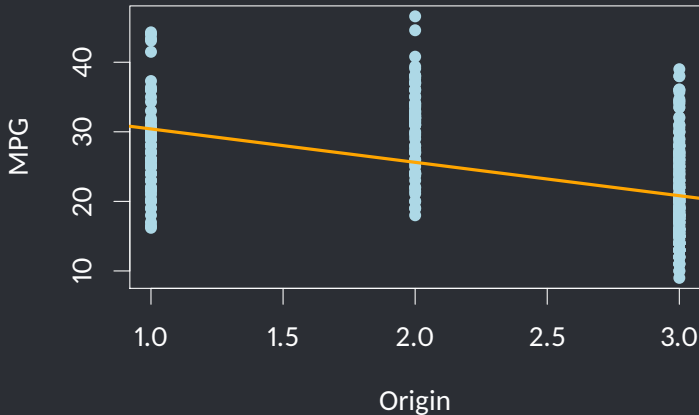
A warning about categorical variables with numeric representations

- In the original dataset, the origin was represented as 1 for US, 2 for EU and 3 for JP.
- We would **not** want to just put these numbers in the regression as numbers, because then regression would treat this as if it were a quantitative variable!

A warning about categorical variables with numeric representations

- In the original dataset, the origin was represented as 1 for US, 2 for EU and 3 for JP.
- We would **not** want to just put these numbers in the regression as numbers, because then regression would treat this as if it were a quantitative variable!
- Even though the representation in the file is numeric, it is still a categorical variable and should be treated as such.

A warning about categorical variables with numeric representations



Statistical significance of a categorical variable

While dealing with categorical variables, we want to look at the significance of the categorical variable as a whole, rather than looking at p -values of individual dummy variables.

Statistical significance of a categorical variable

While dealing with categorical variables, we want to look at the significance of the categorical variable as a whole, rather than looking at p -values of individual dummy variables.

We want to test the **compound null hypothesis**

$$H_0 : \beta_{US} = \beta_{JP} = 0.$$

Statistical significance of a categorical variable

To do this, we look at the ANOVA table; the p -value on the Origin line (2.4×10^{-5}) is the p -value for the compound null hypothesis

$$H_0 : \beta_{US} = \beta_{JP} = 0.$$

```
anova(origin.model)
```

Analysis of Variance Table

Response: MPG

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
HP	1	14433	14433	1096.0	< 2e-16 ***
Weight	1	2392	2392	181.6	< 2e-16 ***
After1975	1	1623	1623	123.2	< 2e-16 ***
Origin	2	288	144	10.9	2.4e-05 ***
Residuals	386	5083	13		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Practical significance of a categorical variable

Since $p < .05$, we can conclude that Origin is a statistically significant predictor of MPG. But is it a *practically* significant predictor?

Practical significance of a categorical variable

Since $p < .05$, we can conclude that Origin is a statistically significant predictor of MPG. But is it a *practically* significant predictor?

To do this, compare R^2 values, or standard error of residuals:

Model	R^2	Residual standard error
Origin not included in model	0.77	3.72
Origin included in model	0.79	3.63

We have to decide if the increased precision is worth the extra complexity in the model.